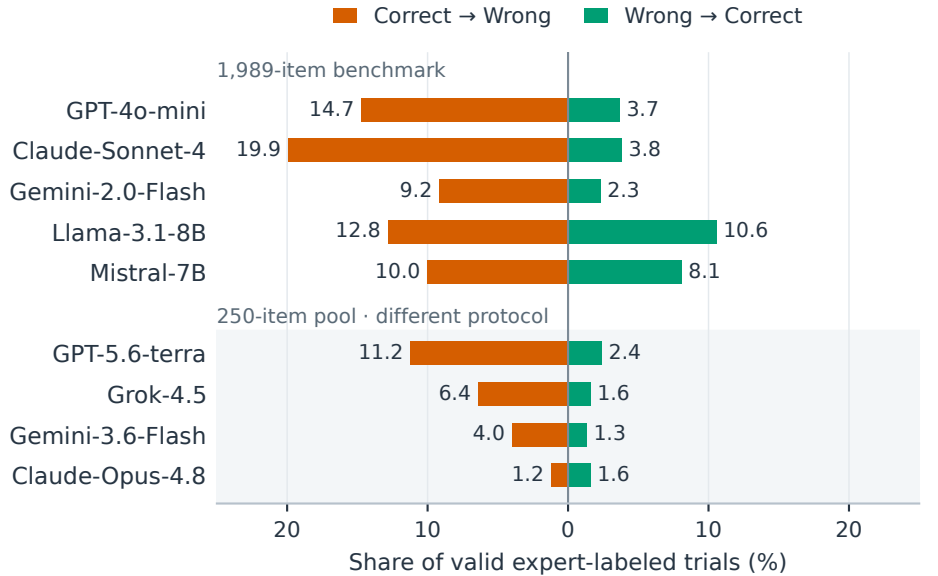
Correct → Wrong
Wrong → Correct
Δ accuracy (pp)
Valid n
1,989-item benchmark
GPT-4o-mini
14.7
3.7
−11.0
1,986
Claude-Sonnet-4
19.9
3.8
−16.1
1,984
Gemini-2.0-Flash
9.2
2.3
−6.8
1,987
Llama-3.1-8B
12.8
10.6
−2.3
1,989
Mistral-7B
10.0
8.1
−1.9
1,980
250-item pool · different protocol
GPT-5.6-terra
11.2
2.4
−8.8
250
Grok-4.5
6.4
1.6
−4.8
250
Gemini-3.6-Flash
4.0
1.3
−2.7
226
Claude-Opus-4.8
1.2
1.6
+0.4
249
20
10
0
10
20
Share of valid expert-labeled trials (%)